

Final presentation

Github version

Adam Wuilmart

August Regnell

Eric Bojs

Ludvig Wärnberg Gerdin

Executive summary



Issues

High computational waste due to difficulty determining feasibility in a optimization problem



Method

Machine learning algorithms



Solutions

SVM

Random Forest

Regression

Deep Learning

Dimensionality
Reduction



Results

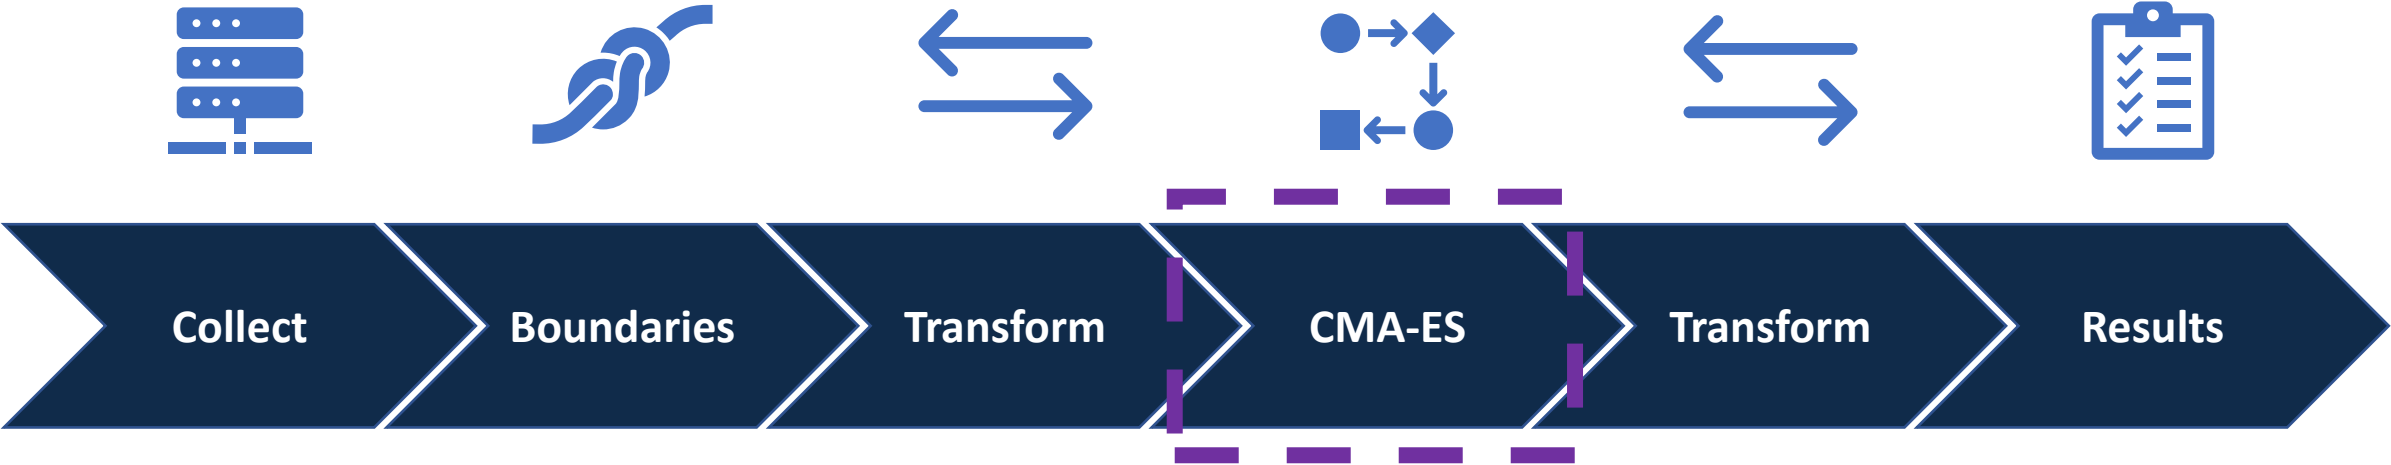
No significant results with current approach due to nature of problem optimization problem



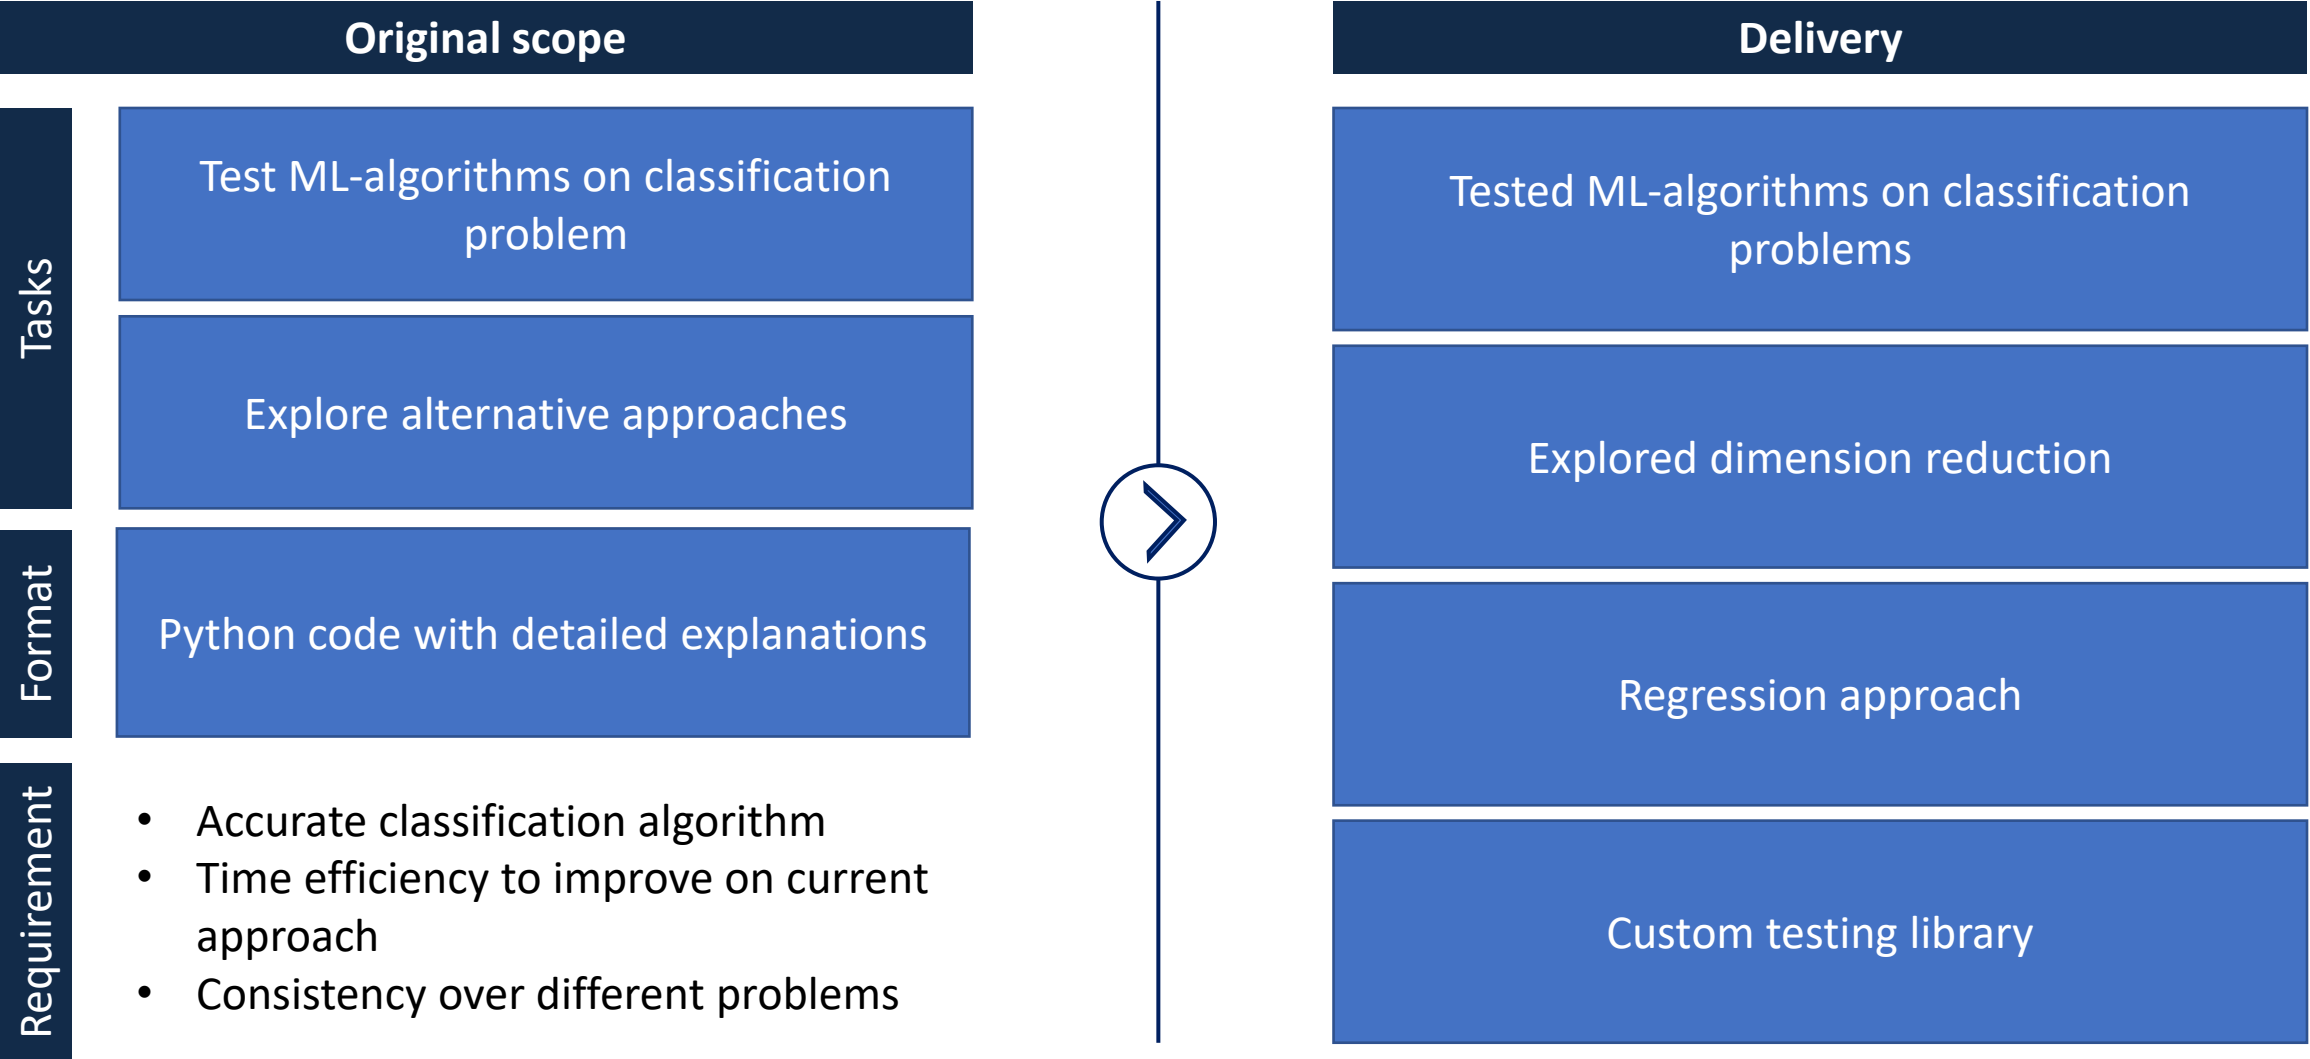
Background

Working with a market technology leader

Client's Value chain: Algorithms to be used in the CMA-ES step



Scope of Work: Additional modules delivered on top of original SOW





Problem

Complex non-linear optimization problem

Background: Client solves complex non-linear optimization problems

Problem formulation


Client finds new trades for a set of banks which will reduce the sum of **CENSORED** required in the network.

$$\begin{aligned} \min f(x) \\ \text{s. t. } x \in D \end{aligned}$$

f is proxy for **total CENSORED** in network
 $x \in R^{4,000}$ is proxy of **new trades** in the network
 D is the **feasible** space, by SIMM & CCP constraints
Both f and D are **non-trivial**

Extremely computer intensive problem to solve


CMA-ES



Each iteration of the evolutionary algorithm **samples 1,024 points** in a region. Next steps **moves towards the optimum** while still being in the feasible region


Data: Data set in use is heavy and highly mobile yielding problems for ML

Data Characteristics



Extremely large and imbalanced dataset

CMA-ES Visualization

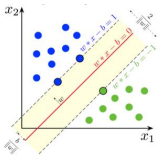
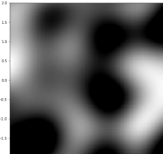
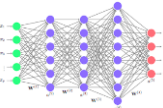
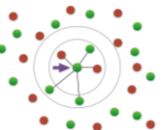
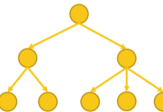




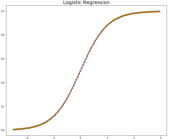

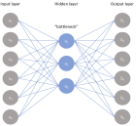
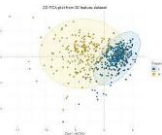
Solution Approach

Classifying points using machine learning

Algorithms: Machine learning algorithms (I/II)

Algorithm	Description	
 <p>A 2D scatter plot with blue and green data points. A red line represents the decision boundary, flanked by yellow shaded regions representing the margin. Dashed lines and mathematical expressions like $w \cdot x + b = 1$ and $w \cdot x + b = -1$ are shown.</p>	Support Vector Machines	Support vector machines is a supervised ML-algorithm which finds the best fitting hyperplane which linearly divides the data
 <p>A grayscale image of the digit '5' on a coordinate grid, representing a function sampled from a Radial Basis Function.</p>	Radial Basis Function Sampler	RBFSampler is a way to approximate the RBF kernels in SVM methods, but with a benefit of being very quick. One combines the approximated kernel with a classifier. The extra speed yields better optimization of parameters
 <p>A diagram of a deep neural network with multiple layers of nodes (neurons) connected by weights.</p>	Deep Learning Networks	Deep learning is part of a broader family of machine learning methods based on artificial neural networks with representation learning. Learning can be supervised, semi-supervised or unsupervised
 <p>A diagram showing a central point surrounded by several other points. Concentric circles represent the search radius for finding the nearest neighbor.</p>	Nearest Neighbor	Nearest Neighbour search, as a form of proximity search, is the optimization problem of finding the point in a given set that is closest, or most similar, to a given point
 <p>A diagram of a decision tree structure with a root node, internal nodes, and leaf nodes.</p>	Random Forest	Random forest classifiers are an ensemble of decision trees, each giving their own classification. The decision made by the ensemble is determined by voting, where the option with most votes from the trees wins

Algorithms: Machine learning algorithms (II/II)

Algorithm	Description
 Logistic Regression	Logistic regression fits a cumulative probability density function using the logit function to the data. This can then be used to classify future samples
 Gradient Boosting Machine	Gradient Boosting Machines build an ensemble of shallow trees, in contrast to Random Forest models that build an ensemble of deep trees. Trees are added <i>sequentially</i> , each new tree learn with respect to current error
 Autoencoder	An autoencoder is a type of artificial neural network used to learn efficient coding of unlabelled data. The encoding is validated and refined by attempting to regenerate the input from the encoding
 Principal Component Analysis	Principal Component Analysis is the process of computing the principal components and using them to perform a change of basis on the data, using only the first few principal components and ignoring the rest

Method: Input from multiple sources yielding standardized testing

Client suggestions




ML-expert suggestions

1. At least twice as many training samples as features
2. 75% train, 25% test
3. Dimensionality reduction infeasible
4. More data, 100k samples is not enough
5. The problem is very difficult



Standardized model tester



We developed a **python library** with:

1. Fixed number of training samples (ca. 8000)
2. Fixed number of test samples (ca. 2000)
3. Fixed lag between training and test samples
4. Hyperparameter tuning
5. Several epochs to get confidence intervals

Output:

variable	average	std	95.0% CI	min	max
weighted accuracy [%]	49.65	1.2	(48.79, 50.51)	47.23	51.73
duration [s]	2.12	0.03	(2.1, 2.14)	2.09	2.18
infeasible_percentage [%]	82.56	2.4	(80.85, 84.28)	77.25	85.45
infeasible_guessed_percentage [%]	35.62	14.46	(25.27, 45.96)	14.84	54.79
feasible_recall [%]	63.8	14.92	(53.12, 74.47)	42.11	85.45
feasible_precision [%]	17.26	2.56	(15.43, 19.09)	13.74	22.65
infeasible_recall [%]	35.5	14.41	(25.19, 45.81)	14.91	54.86
infeasible_precision [%]	82.33	2.47	(80.56, 84.09)	76.92	86
AUC	0.5	0.01	(0.49, 0.51)	0.47	0.52

A hand holding a blue pen points to a bar chart on a document. The chart has several bars with segments in blue, red, and yellow. To the right of the bar chart, there is a line graph with green and red lines. The background is a blurred office setting.

Results

ML-algorithms unfit for problem


Results: RBFSampler + Classifier showing the most promising results

Algorithm	Balanced Accuracy	95% Prediction Interval
Support Vector Machine	50.2%	(49.31%, 51.03%)
RBFSampler + Classifier	51.3%	(50.30%, 52.20%)
Nearest Neighbor	50.3%	(49.51%, 51.01%)
Logistic Regression	50.4%	(49.89%, 50.97%)
Random Forest Classifier	50.4%	(49.75%, 50.58%)
Deeplearning (kNN)	50.2%	(50.04%, 50.30%)
Gradient Boosting Machine (GBM)	50.6%	(49.69%, 51.55%)
Autoencoder + GBM	50.1%	(49.82%, 50.33%)

Support Vector Machine: Slow algorithm with poor results

Model

Support vector machines are supervised ML-algorithms which finds the best fitting hyperplane which linearly divides the data.



Hyper parameters

Kernel function	Transforms hyperspace since the original is not linearly seperable
Regularization parameter	Amount of missclassifications allowed on train data
Kernel coefficient	Specific to kernel chosen



Results

Weighted accuracy: 50.17%

Precision (inf./feasible): 82.42% / 17.63%

Recall (inf./feasible): 46.55% / 53.79%

Speed: 51.69 s




variable	average	std	95.0% CI	min	max
weighted accuracy [%]	50.17	1.2	(49.31, 51.03)	47.82	51.91
duration [s]	51.69	0.13	(51.59, 51.78)	51.56	51.93
infeasible_percentage [%]	82.56	2.4	(80.85, 84.28)	77.25	85.45
infeasible_guessed_percentage [%]	46.5	23.08	(29.99, 63.02)	3.56	79.93
feasible_recall [%]	53.79	22.77	(37.51, 70.08)	23.2	96.14
feasible_precision [%]	17.63	2.35	(15.95, 19.31)	14.66	22.68
infeasible_recall [%]	46.55	23.17	(29.97, 63.12)	3.48	80.51
infeasible_precision [%]	82.42	3.27	(80.08, 84.76)	75.34	85.9
AUC of ROC	0.5	0.02	(0.48, 0.52)	0.46	0.54

Support Vector Machine: Slow algorithm with poor results

Model

Support vector machines are supervised ML-algorithms which finds the best fitting hyperplane which linearly divides the data.



Hyper parameters

Kernel function	Transforms hyperspace since the original is not linearly seperable
Regularization parameter	Amount of missclassifications allowed on train data
Kernel coefficient	Specific to kernel chosen



Results

Weighted accuracy: 50.17%

Precision (inf./feasible): 82.42% / 17.63%

Recall (inf./feasible): 46.55% / 53.79%

Speed: 51.69 s




variable	average	std	95.0% CI	min	max
weighted accuracy [%]	50.17	1.2	(49.31, 51.03)	47.82	51.91
duration [s]	51.69	0.13	(51.59, 51.78)	51.56	51.93
infeasible_percentage [%]	82.56	2.4	(80.85, 84.28)	77.25	85.45
infeasible_guessed_percentage [%]	46.5	23.08	(29.99, 63.02)	3.56	79.93
feasible_recall [%]	53.79	22.77	(37.51, 70.08)	23.2	96.14
feasible_precision [%]	17.63	2.35	(15.95, 19.31)	14.66	22.68
infeasible_recall [%]	46.55	23.17	(29.97, 63.12)	3.48	80.51
infeasible_precision [%]	82.42	3.27	(80.08, 84.76)	75.34	85.9
AUC of ROC	0.5	0.02	(0.48, 0.52)	0.46	0.54

Gradient Boosting Machine: Poor results due to overfitting

Model

Ensemble of shallow trees.
Trees are added *sequentially*, each new tree learn with respect to current error.



Hyper parameters

Maximum tree depth	Adjusts the maximum allowed depth of each tree in ensemble
Number of leaves	Maximum number of leaves in one tree
Regularization parameter	Amount of regularization to apply. That is, to shrink the coefficients.




Results

Weighted accuracy: 50.62%

Precision (inf./feasible): 81.69% / 18.16%

Recall (inf./feasible): 96.88% / 4.36%

Speed: 13.87 s

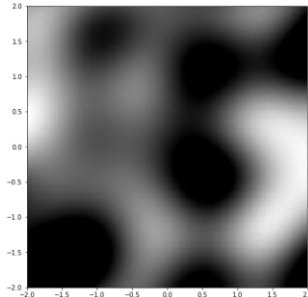


variable	average	std	95.0% PI	min	max
weighted accuracy [%]	50.62	0.75	(49.69, 51.55)	49.93	51.87
duration [s]	13.87	0.61	(13.12, 14.62)	13.25	14.99
infeasible_percentage [%]	81.47	2.67	(78.17, 84.78)	77.25	85.06
infeasible_guessed_percentage [%]	96.62	3.05	(92.83, 100.41)	91.5	99.95
feasible_recall [%]	4.36	4.19	(-0.84, 9.57)	0	11.5
feasible_precision [%]	18.16	11.49	(3.89, 32.42)	0	32.5
infeasible_recall [%]	96.88	2.78	(93.42, 100.34)	92.23	99.94
infeasible_precision [%]	81.69	2.49	(78.6, 84.79)	77.64	85.04

RBFSampler: Fast algorithm with most promising results

Model

The RBFSampler is a way to approximate the RBF kernels in an efficient way. One combines the approximated kernel with a classifier.



Hyper parameters

Length scale (smoothness)

Adjusts the influence of single data points

Number of components

How many Monte Carlo samples that are used per feature

Parameters of the classifier

Loss-function, penalty-type and penalty-strength: depends on the chosen model

Results

Weighted accuracy: 51.25%

Precision (inf./feasible): 83.23% / 18.25%

Recall (inf./feasible): 63.31% / 39.19%

Speed: 3.74 s

3/5

variable	average	std	95.0% CI	min	max
weighted accuracy [%]	51.25	1.32	(50.3, 52.2)	49.01	53.18
duration [s]	3.74	0.81	(3.16, 4.32)	3.21	6.04
infeasible_percentage [%]	82.56	2.4	(80.85, 84.28)	77.25	85.45
infeasible_guessed_percentage [%]	62.89	12.49	(53.95, 71.82)	45.61	84.33
feasible_recall [%]	39.19	13.93	(29.23, 49.16)	17.24	59.44
feasible_precision [%]	18.25	2.17	(16.7, 19.8)	14.72	21.2
infeasible_recall [%]	63.31	12.19	(54.59, 72.03)	46.55	84.62
infeasible_precision [%]	83.23	2.57	(81.39, 85.06)	76.8	86.07


Note: See appendix for remaining algorithms

18

Deeplearning kNNs: Key takeaway from the models results

k-layer Neural Networks

Using multiple layers of connected neurons features are extracted from the input in a similar way to human learning.



Hyper parameters

k	The number of hidden layers in the network
Learning rate	How fast the network learns from every iteration
Optimizer	How and in which direction the learning algorithm moves

Results

Weighted accuracy: 50.17%

Precision (inf./feasible): 82.62% / 21.09%

Recall (inf./feasible): 92.51% / 7.82%

Speed: 1.24 s




variable	average	std	95.0% CI	min	max
weighted accuracy [%]	50.17	0.18	(50.04, 50.3)	50	50.58
duration [s]	1.24	0.18	(1.1, 1.37)	1.05	1.64
infeasible_percentage [%]	82.56	2.4	(80.85, 84.28)	77.25	85.45
infeasible_guessed_percentage [%]	92.45	11.58	(84.17, 100.74)	68.07	100
feasible_recall [%]	7.82	11.86	(-0.66, 16.31)	0	32.92
feasible_precision [%]	21.09	28.63	(0.61, 41.57)	0	100
infeasible_recall [%]	92.51	11.52	(84.27, 100.75)	68.25	100
infeasible_precision [%]	82.62	2.38	(80.92, 84.33)	77.38	85.45
auc of roc	0.49	0.01	(0.48, 0.5)	0.46	0.51

Gradient Boosting Machine: Poor results due to overfitting

Model

Ensemble of shallow trees.
Trees are added *sequentially*, each new tree learn with respect to current error.



Hyper parameters

Maximum tree depth	Adjusts the maximum allowed depth of each tree in ensemble
Number of leaves	Maximum number of leaves in one tree
Regularization parameter	Amount of regularization to apply. That is, to shrink the coefficients.




Results

Weighted accuracy: 50.62%

Precision (inf./feasible): 81.69% / 18.16%

Recall (inf./feasible): 96.88% / 4.36%

Speed: 13.87 s




variable	average	std	95.0% PI	min	max
weighted accuracy [%]	50.62	0.75	(49.69, 51.55)	49.93	51.87
duration [s]	13.87	0.61	(13.12, 14.62)	13.25	14.99
infeasible_percentage [%]	81.47	2.67	(78.17, 84.78)	77.25	85.06
infeasible_guessed_percentage [%]	96.62	3.05	(92.83, 100.41)	91.5	99.95
feasible_recall [%]	4.36	4.19	(-0.84, 9.57)	0	11.5
feasible_precision [%]	18.16	11.49	(3.89, 32.42)	0	32.5
infeasible_recall [%]	96.88	2.78	(93.42, 100.34)	92.23	99.94
infeasible_precision [%]	81.69	2.49	(78.6, 84.79)	77.64	85.04

Nearest Neighbor: Key takeaway from the models results

Nearest Neighbor

Classifying new points be performing a plurality vote of its k nearest neighbors.



Hyper parameters

k	The number of neighbors considered.
Distance metric	How the distances should be calculated, e.g., l2-norm



Results

Weighted accuracy: 50.26%

Precision (inf./feasible): 82.48% / 17.65%

Recall (inf./feasible): 50.58% / 49.94%

Speed: 1.11 s




variable	average	std	95.0% CI	min	max
weighted accuracy [%]	50.26	1.04	(49.51, 51.01)	48.35	52.12
duration [s]	1.11	0.2	(0.96, 1.25)	0.79	1.43
infeasible_percentage [%]	82.56	2.4	(80.85, 84.28)	77.25	85.45
infeasible_guessed_percentage [%]	50.51	36.55	(24.36, 76.66)	7.37	96.29
feasible_recall [%]	49.94	35.84	(24.3, 75.57)	2.65	92.75
feasible_precision [%]	17.65	2.85	(15.62, 19.69)	11.84	22.6
infeasible_recall [%]	50.58	36.7	(24.33, 76.84)	7.4	96.08
infeasible_precision [%]	82.48	2.7	(80.55, 84.41)	76.81	85.85
auc of roc	0.5	0.01	(0.5, 0.51)	0.48	0.52

Logistic Regression: Key takeaway from the models results

Logistic Regression

Logistic regression fits a cumulative probability density function using the logit function to the data. This can then be used to classify future samples.



Hyper parameters

Penalty	The loss function assigned to missclassifications
Regularization parameter	Amount of missclassifications allowed on train data
Fit interception	Weather or not we should assume that $\beta_0=0$ or not

Results

Weighted accuracy: 50.43%

Precision (inf./feasible): 82.79% / 17.65%

Recall (inf./feasible): 36.19% / 64.68%

Speed: 496.05 s




variable	average	std	95.0% CI	min	max
weighted accuracy [%]	50.43	0.75	(49.89, 50.97)	49.02	51.65
duration [s]	496.05	182.34	(365.61, 626.49)	286.23	707.68
infeasible_percentage [%]	82.56	2.4	(80.85, 84.28)	77.25	85.45
infeasible_guessed_percentage [%]	36.05	9.6	(29.18, 42.92)	22.61	49.71
feasible_recall [%]	64.68	8.87	(58.33, 71.02)	51.31	77.66
feasible_precision [%]	17.65	2.3	(16.01, 19.3)	15.05	22.48
infeasible_recall [%]	36.19	9.76	(29.2, 43.17)	22.5	50.06
infeasible_precision [%]	82.79	2.69	(80.87, 84.71)	76.52	86.3
AUC of ROC	0.49	0.01	(0.48, 0.5)	0.47	0.51

Random Forest:

Random forest

Random forest classifiers ensembles of decision trees, each giving a classification. Decision made by the ensemble is determined by voting, where the option with most votes from the trees wins



Hyper parameters

Estimators	Number of trees used in the forest
Criterion	How each split is determined
Max depth	Maximum number of levels allowed in the tree



Results

Weighted accuracy: 50.01%

Precision (inf./feasible): 82.58% / 3.32%

Recall (inf./feasible): 98.2% / 1.9%

Speed: 7 minutes




The model is slow and gives poor accuracy. There will be a high need for hyper-parameter tuning in the case that the model works to prevent it from overfitting the data. This will significantly vary from case to case. Thereby, we do not recommend the use of this model for the problem

Autoencoder + GBM: Key takeaway from the models results

Autoencoder

Neural network used to denoise data. The performance measured by squared cell-wise difference between denoised and original data.



Hyper parameters

Number of layers	Adjust the complexity of the model, I.e. how well-fitted to sample
Learning rate	Distance to jump in each iteration
Activation function	Transformation of output from each layer

Results

Weighted accuracy: 50.08%

Precision (inf./feasible): 82.67% / 12.08%

Recall (inf./feasible): 92.68% / 7.48%

Speed: 198.17 s



variable	average	std	95.0% PI	min	max
weighted accuracy [%]	50.08	0.36	(49.82, 50.33)	49.69	51.08
duration [s]	198.17	20.73	(183.34, 213.0)	170.08	223.45
infeasible_percentage [%]	82.56	2.4	(80.85, 84.28)	77.25	85.45
infeasible_guessed_percentage [%]	92.64	20.25	(78.15, 107.12)	31.93	100
feasible_recall [%]	7.48	20.77	(-7.39, 22.34)	0	69.74
feasible_precision [%]	12.08	11.07	(4.16, 20.0)	0	28.57
infeasible_recall [%]	92.68	20.1	(78.3, 107.05)	32.43	100
infeasible_precision [%]	82.67	2.14	(81.14, 84.2)	78.44	85.43


A man with short dark hair is seen from the side, looking at a large bulletin board. The board is covered with numerous photographs, documents, and a large map. The man is wearing a white shirt and a dark jacket. The background is slightly blurred, focusing attention on the man and the bulletin board.

Underlying Causes


Data, dimensionality & mobility

Outcome analysis: Mobility of CMA-ES causing difficulties for ML


3D PCA on all iterations



2D PCA on 7 iterations




Distance between iterations



Jumps between iterations poses a very difficult challenge for ML-algorithms due to extrapolation

Outcome analysis: Infeasibility of feature reduction due to boundaries


Stability of influential features



If parameter β_i is influential, we would expect it to show up in both sets. However, this was **not found**

Principal Component Analysis

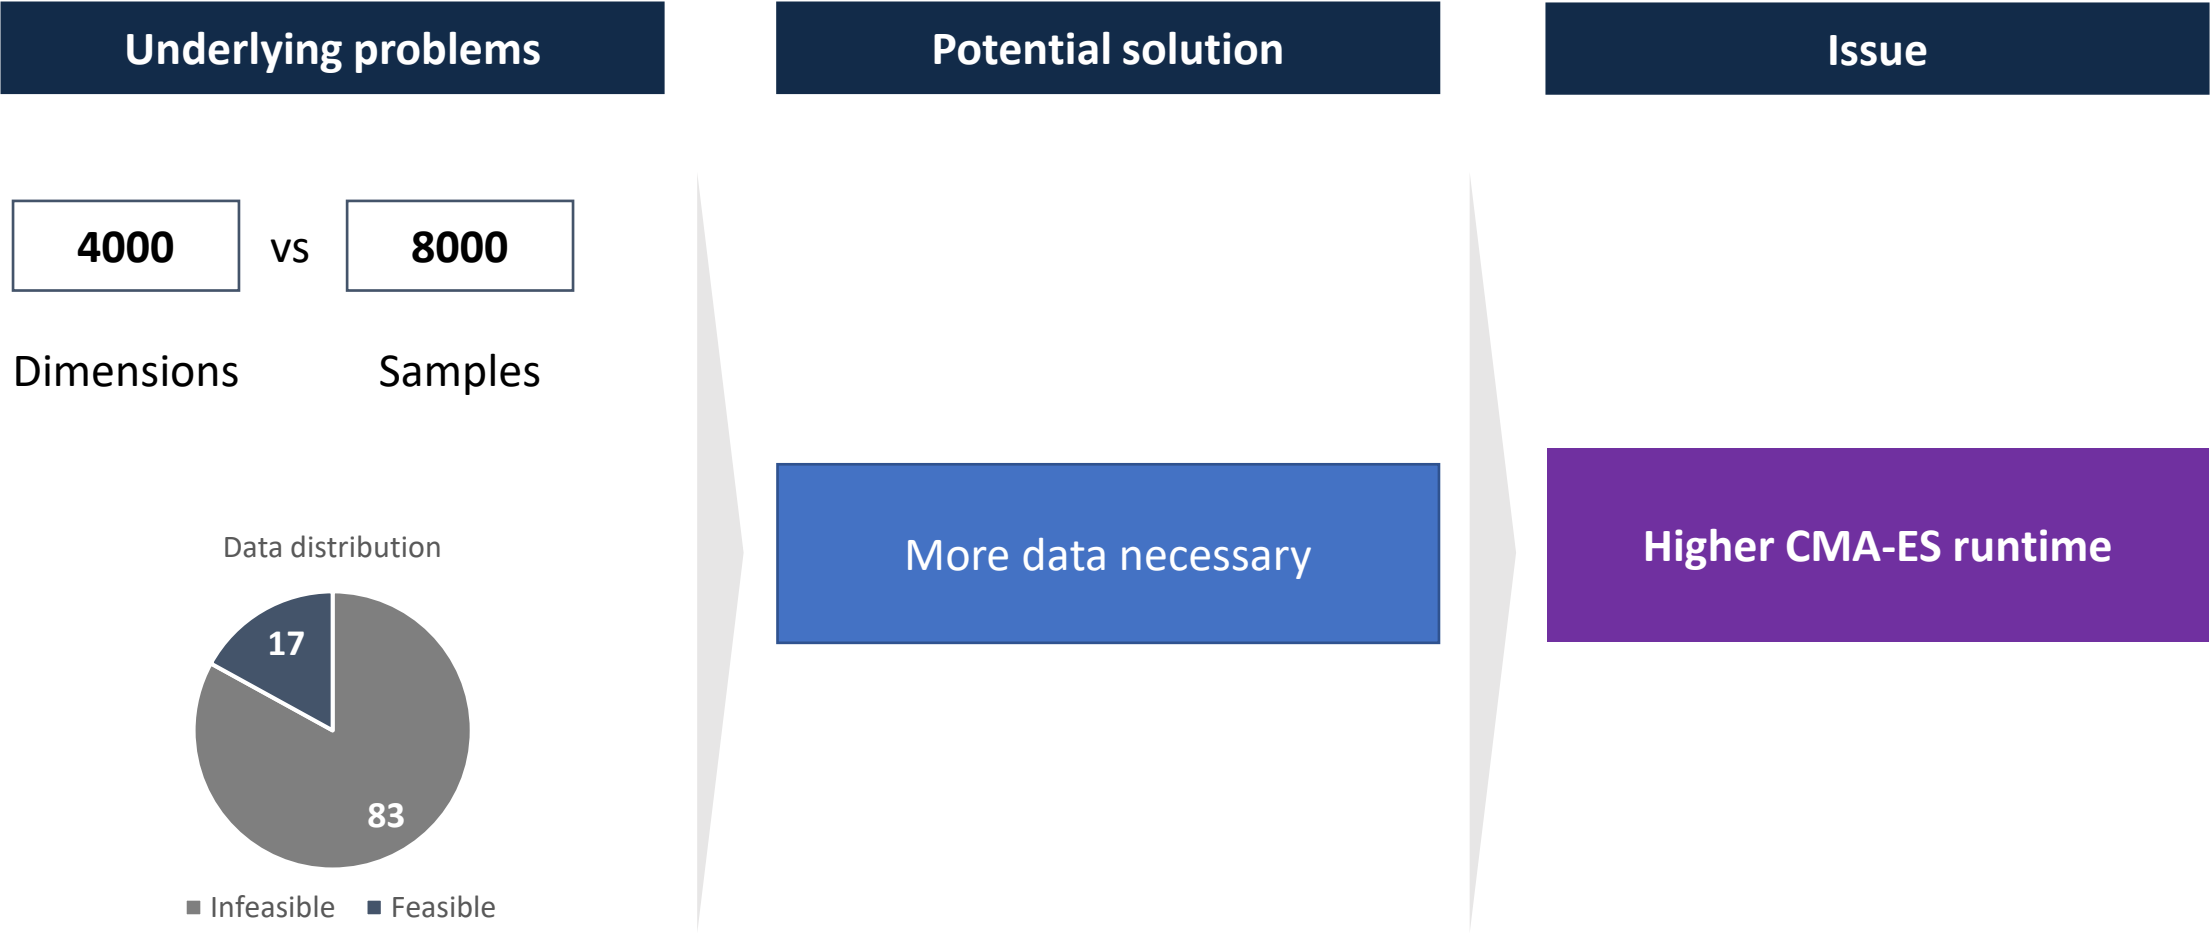
We perform PCA on the in-/feasible solutions separately and see if there is redundant information



Feature reduction could **not be done**


ML-algorithms have trouble training on few samples in the vicinity of test set which yields insufficient information

Outcome analysis: Insufficient data for algorithms to work with




Outcome analysis: Results due to mobility, information and data size

CMA-ES mobility



Information density



Data size

4000

vs

8000

Dimensions

Samples

Insignificant results from ML algorithms

Next steps: Two potential keys to unlock the problem



Transforms and adaptation
possible with algorithm
understanding



Train on larger datasets or
change algorithm approach

